

1. Introduction
 - a. About me
 - i. First time being an instructor of record for a class
 - ii. If you have any comments or constructive feedback, let me know!
 - b. About this class (why computer architecture?)
 - i. This is a required class for both CS and CSE
 - ii. “Core material” for most CS graduate programs
 - iii. My goal is to get you to remember some of this knowledge to use in industry or your next computer architecture class
 - iv. My *hope* is that I get some of you interested in computer architecture
2. About the class itself
 - a. Discussions on second half of Wednesdays by our TA, Yuan
 - b. Three quizzes (20%), one midterm (20%), and one final (30%)
 - i. All of these will be on Thursdays, which are in a different room!
 - ii. First quiz on Thursday, will get back to you before the drop deadline
 - c. Four to five labs (30%, still debating on number)
 - i. Mostly in Logisim, one might be in C++
 - ii. First assignment will get you used to Logisim
 - iii. Can turn in up to 48 hours late for non-linear penalty
 - d. Regrades
 - i. One week from return of assignment
 - ii. Office hours only
 - e. Websites (there’s a lot of them)
 - i. Github is where all the course materials get posted
 - ii. Canvas is for submitting assignments
 - iii. Gradescope is for returning your tests
 - iv. Piazza is for discussion
 - f. Textbook
 - i. Reading schedule on Github
 - ii. *Digital McLogic Design* provided on Canvas for some of the digital design stuff
 - iii. Otherwise, *Computer Organization and Design* for everything else
 1. 9th or 10th edition is fine
3. Computer architecture
 - a. Abstraction layers
 - i. Software
 1. Applications, OS
 2. Could divide the OS further into file system, drivers, kernels...
 - ii. Hardware
 1. Hardware devices, gates, wires, transistors, electrons
 2. Again, could subdivide further
 - iii. Computer architecture
 1. The “hardware-software interface” according to Hennessey and Patterson of Stanford and UCB

- iv. Computer architecture is a broad topic
 - 1. You're not going to know everything after this class
 - 2. Get a general overview
 - 3. Some of the material could be useful in industry, say, a SWE position
 - a. Example: designing an application's working set to fit in L2 cache
 - v. What we're going to talk about
 - 1. Digital design – basic building blocks
 - 2. Design parts of a CPU and memory
 - 3. Learn about busses and about memory
 - 4. Do this in a different order than the book
 - a. Memory is important
 - b. Bussing is also important, but doesn't lend itself to good homework
 - b. Definitions
 - i. Computer architecture
 - 1. Attributes of a system visible to the programmer
 - 2. Those which impact the logical execution of a program
 - ii. Computer organization
 - 1. Operations, units
 - 2. Their interconnections that realize the architectural specifications
 - c. Tasks of a computer
 - i. Transfer data between external devices
 - 1. Keyboard to monitor
 - 2. Microphone to speaker, so on
 - ii. Storage device
 - 1. Network to memory
 - iii. Data processing
 - 1. Internal or external source and destination
 - iv. Control external devices
 - d. Parts of a computer
 - i. I/O
 - 1. Mouse, keyboard main examples we think of
 - 2. Other peripherals count too (like speakers)
 - ii. Main memory: RAM, caches
 - iii. System bus
 - iv. CPU
 - 1. Registers, store values
 - 2. ALU, perform operations
 - 3. Internal bus, transfer data
 - 4. Control unit
 - a. Sequencing logic, where to go next
 - b. Control unit registers, decoders
 - c. Control memory
4. History of computing
- a. Mechanical
 - i. 1801 – Joseph-Marie Jacquard's loom
 - ii. 1842 – Charles Babbage
 - 1. Difference engine to compute polynomials
 - 2. Analytical engine, like modern machines

- b. Electromechanical
 - i. 1936 – Alan Turing “universal computing machine”
 - ii. Data and program on single tape
 - c. Electronic
 - i. First generation – vacuum tubes
 - 1. 1945 – John Von Neumann working on the Manhattan Project
 - a. With team, designs architecture that is used by nearly every machine today
 - b. Von Neumann architecture – get to this later
 - 2. 1946 – Electronic Numerical Integrator Analyzer and Computer, or ENIAC
 - a. Used to calculate trajectories for bombing runs
 - b. Not very good – 20% made it within 1000'
 - c. Huge, 80' long and 8.5' high
 - d. Difficult to program, using patch boards
 - ii. Second generation – transistors
 - 1. 1959 – memory via small ferro-magnetic donuts with wires running through them
 - 2. Source of the term “core memory”
 - 3. 1K memory was the size of a shoe box
 - iii. Third generation – planar transistors
 - 1. 1964 – transistors on pieces of silicon
 - 2. Printed circuit boards (PCBs)
 - iv. Fourth generation – CPU on one chip
 - 1. 1971 – use very-large-scale-integration (VLSI) to accomplish
 - v. Fifth generation – now(?)
 - 1. Lines are fuzzy
 - 2. Standard chips, programmable logic devices (PLAs)
 - 3. Field programmable gate array (FPGAs)
 - 4. Application specific integrated circuits (ASICs)
 - a. Today, commonly heard with crypto mining
5. Designing for performance
- a. Increase CPU performance
 - i. Increase clock frequency
 - ii. Increase size and speed of caches on CPU
 - iii. More parallelism
 - 1. Pipelining and other instruction-level parallelism techniques
 - 2. Branch prediction
 - 3. Speculative execution
 - iv. Improve interface between RAM and CPU
 - b. Power wall
 - i. Want to increase clock frequency
 - ii. $P = cfv^2$ (p = power, c = capacitance, f = frequency, v = voltage)
 - iii. Lower voltage to increase frequency, effective because of the squared term
 - iv. Need a certain voltage to differentiate 0 and 1, we're at that point
 - v. Can't increase power further without burning chips
 - vi. Hit this power wall around 2006
 - c. Gordon Moore's law, 1965
 - i. Every 18 months, number of transistors on a chip will double
 - ii. Still holding (roughly), but dark silicon problem
 - iii. Can't power all the transistors at once without burning the chip

- d. (Robert) Dennard scaling, 1974
 - i. As transistors get smaller, power density stays constant
 - ii. Make a transistor smaller, uses less power
 - iii. Broke down around 2006